

```
In [150]: 1 import numpy as np
2 import pandas as pd
3 import seaborn as sns
4 from tqdm.notebook import tqdm
5 import matplotlib.pyplot as plt
6
7 import torch
8 import torch.nn as nn
9 import torch.optim as optim
10 from torch.utils.data import Dataset, DataLoader, WeightedRandomSampler
11
12 from sklearn.preprocessing import MinMaxScaler
13 from sklearn.model_selection import train_test_split
14 from sklearn.metrics import confusion_matrix, classification_report
15
16 import sklearn as skl
17 import skorch as skr
```

```
In [155]: 1 print('pands', pd.__version__)
2 print('numpy', np.__version__)
3 print('seaborn', sns.__version__)
4 print('torch', torch.__version__)
5 print('skorch', skr.__version__)
```

```
pands 1.4.2
numpy 1.21.5
seaborn 0.11.2
torch 1.13.1+cpu
skorch 0.12.1
```

```
In [2]: 1 df = pd.read_excel('ctgdata.xlsx')
2 df.head()
```

Out[2]:

	Unnamed: 0	b	e	LBE	LB	AC	FM	UC	ASTV	MSTV	...	C	D	E	AD	DE	LD	FS	SUSP	CLASS	NSP
0	0	240	357	120	120	0	0	0	73	0.5	...	0	0	0	0	0	0	1	0	9	2
1	1	5	632	132	132	4	0	4	17	2.1	...	0	0	0	1	0	0	0	0	6	1
2	2	177	779	133	133	2	0	5	16	2.1	...	0	0	0	1	0	0	0	0	6	1
3	3	411	1192	134	134	2	0	6	16	2.4	...	0	0	0	1	0	0	0	0	6	1
4	4	533	1147	132	132	4	0	5	16	2.4	...	0	0	0	0	0	0	0	0	2	1


5 rows × 37 columns

Data exploration

I know from the initial data exploration done during the group coursework that this dataset has been preprocessed. However during the group cw we decided against any feature extraction/transformation because it was not necessary. Some research has suggested feature extraction increases accuracy. Thus, features with low correlations and variances will be removed.

```
In [3]: 1 #class distribution
2 sns.countplot(x = 'NSP', data=df)
```

Out[3]: <AxesSubplot:xlabel='NSP', ylabel='count'>



```
In [5]: , 1ALTV', 'MLTV', 'Width', 'Min', 'Max', 'Nmax', 'Nzeros', 'Mode', 'Mean', 'Median', 'Variance', 'Tendency', 'SUSP', 'CLASS'
```

UCI says there are only 23 attributes, but this shows 35. According to UCI feature are: LB, AC, FM, UC, DL, DS, DP, ASTV, MSTV, ALTV, MLTV, Width, Min, Max, NMax, Nzeros, Mode, Mean, Median, Variance, Tendency, Class, NSP. I'll keep SUSP because it appears to have a stornger corr than most. All others will be dropped

```
In [4]: 1 #reduced to 23 attributes.
2 df.head()
```

Out[4]:

	Unnamed: 0	b	e	LBE	LB	AC	FM	UC	ASTV	MSTV	...	C	D	E	AD	DE	LD	FS	SUSP	CLASS	NSP
0	0	240	357	120	120	0	0	0	73	0.5	...	0	0	0	0	0	0	1	0	9	2
1	1	5	632	132	132	4	0	4	17	2.1	...	0	0	0	1	0	0	0	0	6	1
2	2	177	779	133	133	2	0	5	16	2.1	...	0	0	0	1	0	0	0	0	6	1
3	3	411	1192	134	134	2	0	6	16	2.4	...	0	0	0	1	0	0	0	0	6	1
4	4	533	1147	132	132	4	0	5	16	2.4	...	0	0	0	0	0	0	0	0	2	1

5 rows × 37 columns

the data is imbalanced, which was known. SMOTE did little to improve this during prelimbary scikit phase. Will feature extraction improve? /

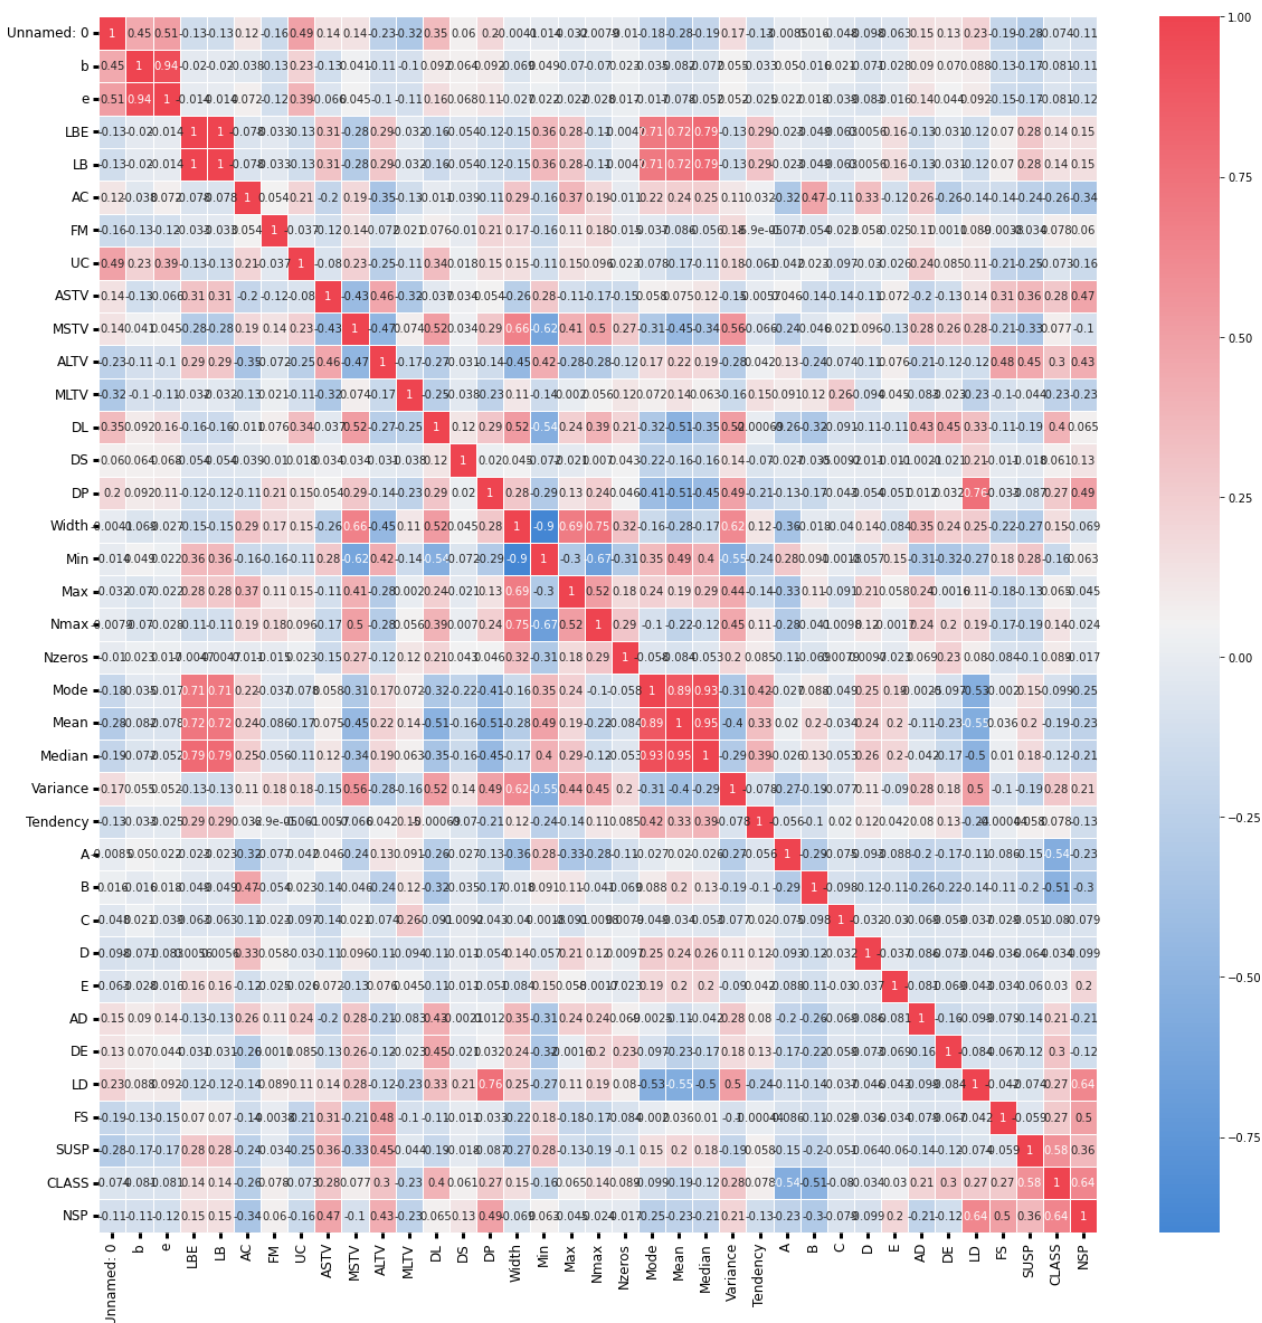
```

In [5]: 1 # borrowed from: https://www.kaggle.com/code/christopherwsmith/fetal-health-a-quick-guide-to-high-accuracy
2 def Plotter(plot, x_label, y_label, x_rot=None, y_rot=None, fontsize=12, fontweight=None, legend=None, save=False, save_name=None):
3     """
4     Helper function to make a quick consistent plot with few easy changes for aesthetics.
5     Input:
6     plot: sns or matplotlib plotting function
7     x_label: x_label as string
8     y_label: y_label as string
9     x_rot: x-tick rotation, default=None, can be int 0-360
10    y_rot: y-tick rotation, default=None, can be int 0-360
11    fontsize: size of plot font on axis, default=12, can be int/float
12    fontweight: Adding character to font, default=None, can be 'bold'
13    legend: Choice of including legend, default=None, bool, True:False
14    save: Saves image output, default=False, bool
15    save_name: Name of output image file as .png. Requires Save to be True.
16              default=None, string: 'Insert Name.png'
17    Output: A customized plot based on given parameters and an output file
18    """
19
20    #Ticks
21    ax.tick_params(direction='out', length=5, width=3, colors='k',
22                  grid_color='k', grid_alpha=1, grid_linewidth=2)
23    plt.xticks(fontsize=fontsize, fontweight=fontweight, rotation=x_rot)
24    plt.yticks(fontsize=fontsize, fontweight=fontweight, rotation=y_rot)
25
26    #Legend
27    if legend==None:
28        pass
29    elif legend==True:
30
31        plt.legend()
32        ax.legend()
33        pass
34    else:
35        ax.legend().remove()
36
37    #Labels
38    plt.xlabel(x_label, fontsize=fontsize, fontweight=fontweight, color='k')
39    plt.ylabel(y_label, fontsize=fontsize, fontweight=fontweight, color='k')
40
41    #Removing Spines and setting up remaining, preset prior to use.
42    ax.spines['top'].set_color(None)
43    ax.spines['right'].set_color(None)
44    ax.spines['bottom'].set_color('k')
45    ax.spines['bottom'].set_linewidth(3)
46    ax.spines['left'].set_color('k')
47    ax.spines['left'].set_linewidth(3)
48
49    if save==True:
50        plt.savefig(save_name)

```

```
In [6]: 1 fig, ax=plt.subplots(figsize=(20,20))#Required outside of function. This needs to be activated first when plotting in
2 cmap = sns.diverging_palette(250, 10, s=80, l=55, n=9, as_cmap=True)
3 plot=sns.heatmap(df.corr(),annot=True, cmap=cmap, linewidths=1)
4 Plotter(plot, None, None, 90,legend=False, save=True, save_name='Corr.png')
```

No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.



Shades of red are more correlated than blue. Looking at the NSP col/row it appears that LB, DS, DP, ASTV, ALTV, Variance, SUSP, CLASS have the best correlation.

```
In [7]: 1 # Using KBest Algo with f_classif to perform ANOVA which:
2 #determines the degree of linear dependency between the target variable and features.
3 from sklearn.feature_selection import SelectKBest #Feature Selector
4 from sklearn.feature_selection import f_classif #ANOVA
```


```

In [8]: 1 #Feature Selection
2 X=df.drop(['NSP'], axis=1)
3 Y=df['NSP']
4 bestfeatures = SelectKBest(score_func=f_classif, k='all')
5 fit = bestfeatures.fit(X,Y)
6 dfscores = pd.DataFrame(fit.scores_)
7 dfcolumns = pd.DataFrame(X.columns)
8 #concat two dataframes for better visualization
9 featureScores = pd.concat([dfcolumns,dfscores],axis=1)
10 featureScores.columns = ['Feature','Score'] #naming the dataframe columns
11
12 #Visualize the feature scores
13 fig, ax=plt.subplots(figsize=(7,7))
14 plot=sns.barplot(data=featureScores, x='Score', y='Feature', palette='viridis',linewidth=0.5, saturation=2, orient='h')
15 Plotter(plot, 'Score', 'Feature', legend=False, save=True, save_name='Feature Importance.png')#Plotter function for aes
16 plot

```

No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.

Out[8]: <AxesSubplot:xlabel='Score', ylabel='Feature'>



SUSP isn't listed in the official attribute information, but it was correlated and now shows it has the highest linear dependency. I'm going to exclude it because it's not listed and could be an outlier. 250 looks to be a good cut off point for feature selection.

```

In [41]: 1 #Selection method
2 selection=featureScores[featureScores['Score']>=200]#Selects features that scored more than 200
3 selection=list(selection['Feature'])#Generates the features into a list
4 selection.append('NSP')#Adding the Level string to be used to make new data frame
5 df_feat=df[selection] #New dataframe with selected features
6 df_feat = df_feat.drop(columns=['SUSP', 'FS', 'LD', 'E'])
7 df_feat.head() #Lets take a look at the first 5

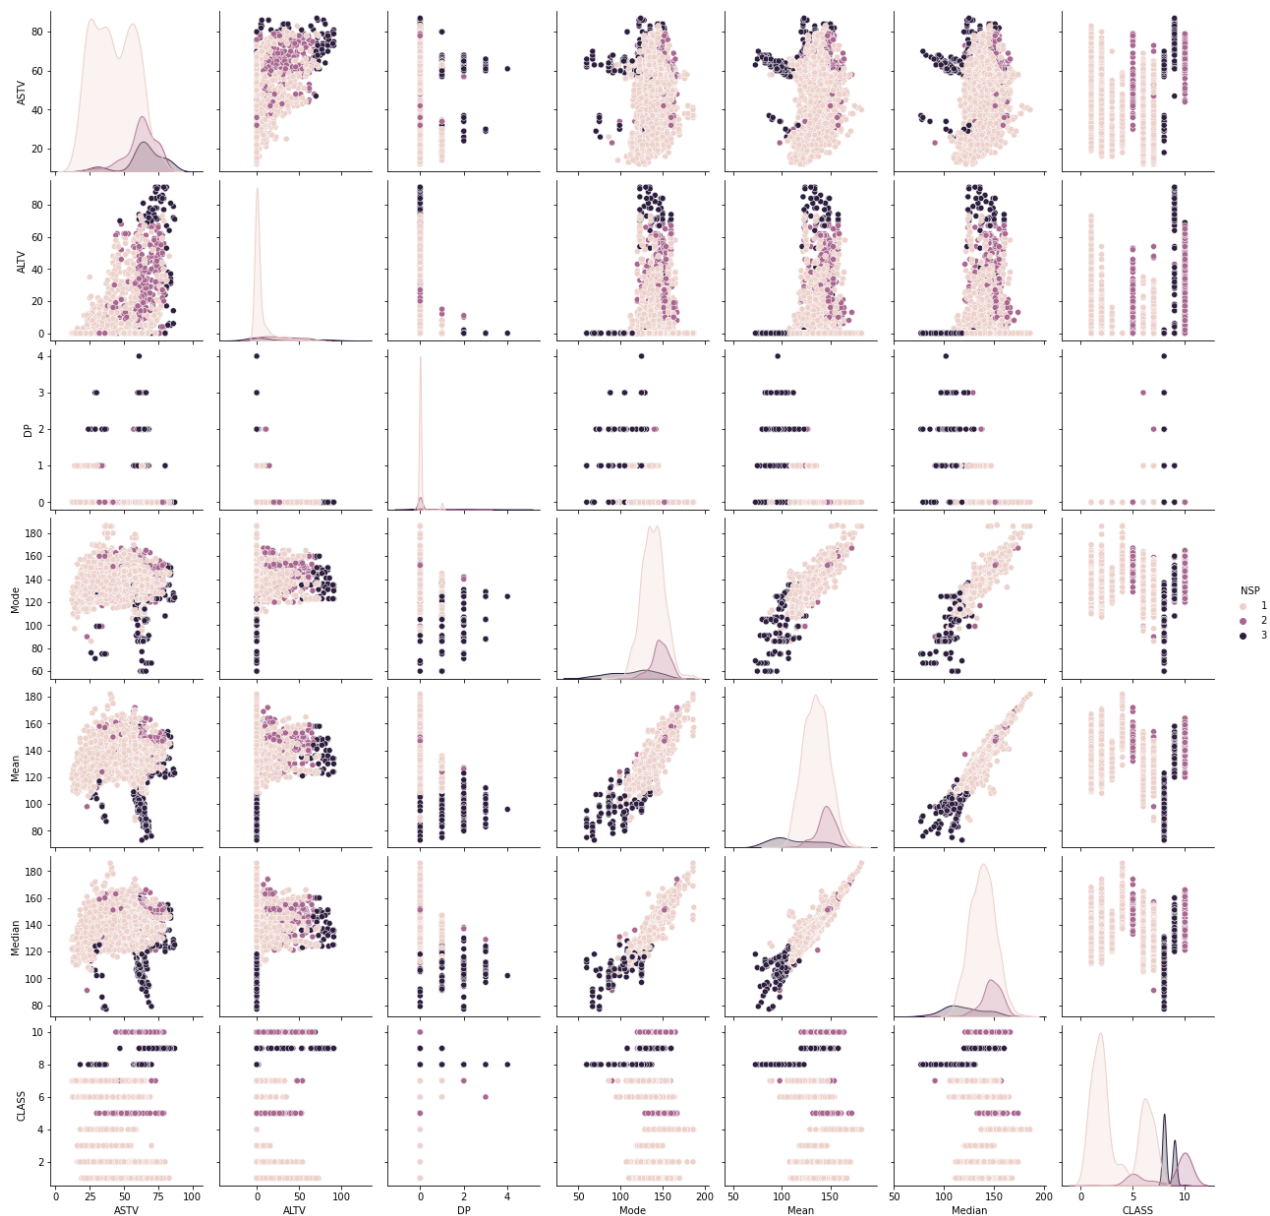
```

Out[41]:

	ASTV	ALTV	DP	Mode	Mean	Median	CLASS	NSP
0	73	43	0	120	137	121	9	2
1	17	0	0	141	136	140	6	1
2	16	0	0	141	135	138	6	1
3	16	0	0	137	134	137	6	1
4	16	0	0	137	136	138	2	1

```
In [42]: 1 sns.pairplot(df_feat, hue='NSP')
2
```

Out[42]: <seaborn.axisgrid.PairGrid at 0x18c53131040>



IEEE paper mentions 7 features so this seems to be a good choice!

Classes 2 and 3 are difficult to distinguish here.

Splitting, sclaing, encoding

```
In [43]: 1 # make things simple
2 data = df_feat
```

```
In [44]: 1 # Encoding the output. Labels need to go from 0-2 in order to work with tensor
2 # 0 = Normal, 1 = Suspect, 2 = Pathologic
3 # borrowed from https://towardsdatascience.com/pytorch-tabular-multiclass-classification-9f8211a123ab
4
5 class2idx = {
6     1:0,
7     2:1,
8     3:2
9 }
10
11 idx2class = {v: k for k, v in class2idx.items()}
12 data['NSP'].replace(class2idx, inplace=True)
```

```
In [97]: 1 #Create inputs and targets.
2 X = data.iloc[:, 0:-1]
3 y = data.iloc[:, -1]
4
```

```
In [98]: 1 print(X.head())
```

	ASTV	ALTV	DP	Mode	Mean	Median	CLASS
0	73	43	0	120	137	121	9
1	17	0	0	141	136	140	6
2	16	0	0	141	135	138	6
3	16	0	0	137	134	137	6
4	16	0	0	137	136	138	2

```
In [99]: 1 print(y.head())
```

```
0    1
1    0
2    0
3    0
4    0
Name: NSP, dtype: int64
```

```
In [100]: 1 # Create split into train, val, test
2 # Split into train+val and test
3 # Stratify is being used to have an equal distribution of output classes sets
4 # test_size is .2, as mentioned in paper
5
6 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, stratify=y, random_state=42)
```

```
In [101]: 1 # Normalize the input. Neural networks need a range of 0,1
2 # Use MinMaxScaler to transform features
3 scaler = MinMaxScaler()
4
5 X_train = scaler.fit_transform(X_train)
6 X_test = scaler.transform(X_test)
7
8 # convert inputs and outputs in numpy arrays
9 X_train, y_train = np.array(X_train), np.array(y_train)
10 X_test, y_test = np.array(X_test), np.array(y_test)
```

```
In [102]: 1 X_test.shape
```

```
Out[102]: (426, 7)
```

X_val and X_test, .transform was used because the validation and test sets should be scaled with the same parameters as the train set to avoid data leakage. fit_transform calculates scaling values and applies. .transform only applies the calculated values.

Cross-validation will be done in the model building phase

Neural Network

Model parameters

```
In [103]: 1 # create tensors
          2 X_train = torch.tensor(X_train)
          3 X_test = torch.tensor(X_test)
          4
          5
          6 y_test = torch.tensor(y_test)
          7 y_train = torch.tensor(y_train)
          8
```

```
In [104]: 1 print(f"Datatypes of training data: X: {X_test.dtype}, y: {y_train.dtype} ")
```

Datatypes of training data: X: torch.float64, y: torch.int64

```
In [105]: 1
          2 from sklearn.model_selection import cross_val_score
          3
```

```
In [106]: 1 #implementing baisc Log regression with default as test?
          2 from sklearn import linear_model
          3 from sklearn.linear_model import LogisticRegression
          4
          5 logistic_regression = linear_model.LogisticRegression()
          6 logistic_regression_mod = logistic_regression.fit(X_train, y_train)
          7 print(f"Baseline Logistic Regression: {round(logistic_regression_mod.score(X_test, y_test), 3)}")
          8
          9 pred_logistic_regression = logistic_regression_mod.predict(X_test)
```

Baseline Logistic Regression: 0.913

Define architecture

Two hidden layers, because of the universal approximation theorem. Input size is 7, output size is 3 classes

```
In [111]: 1 import torch.nn.functional as F
          2 import torch.nn as nn
          3
          4 class ctgClassifier(nn.Module):
          5     def __init__(self, dropout=0.5, weight_constraint=1.0):
          6         super(ctgClassifier, self).__init__()
          7         self.dropout = nn.Dropout(dropout)
          8
          9         self.layer_1 = nn.Linear(7, 128)
          10        self.layer_2 = nn.Linear(128, 64)
          11        self.layer_out = nn.Linear(64, 3)
          12
          13
          14        def forward(self, x):
          15            x = F.relu(self.layer_1(x))
          16            x = self.dropout(x)
          17            x = F.relu(self.layer_2(x))
          18            x = self.dropout(x)
          19            return x
```

The forward pass of the neural network takes an input tensor x and applies the fully connected layers and activation functions defined. The `F.relu()` function applies the ctivation function to the output of each fully connected layer. The `self.dropout(x)` applies dropout regularization to the output of the first and second hidden layers. Finally, the function returns the output tensor x .

In fact, there is a theoretical finding by Lippmann in the 1987 paper "An introduction to computing with neural nets" that shows that an MLP with two hidden layers is sufficient for creating classification regions of any desired shape

Specifically, the universal approximation theorem states that a feedforward network with a linear output layer and at least one hidden layer with any "squashing" activation function (such as the logistic sigmoid activation function) can approximate any Borel measurable function from one finite-dimensional space to another with any desired non-zero amount of error, provided that the network is given enough hidden units.

— Page 198, Deep Learning, 2016


```
In [112]: ▶ from torch.nn import EarlyStopping
from torch.callbacks import EarlyStopping
early_stopping = EarlyStopping(monitor='valid_loss', patience = 10, threshold = 0.0001, threshold_mode='rel', lower_is_better=True)
```

```
In [113]: ▶ 1 #Multi-Layer Perceptron classifier.
2
3 from skorch import NeuralNetClassifier
4
5 net = NeuralNetClassifier(
6     ctgClassifier,
7     lr=0.1,
8     criterion = torch.nn.modules.loss.CrossEntropyLoss,
9     optimizer=torch.optim.Adam,
10    callbacks=[early_stopping]
11 )
12
```

details for tuning: <https://machinelearningmastery.com/how-to-grid-search-hyperparameters-for-pytorch-models/> (<https://machinelearningmastery.com/how-to-grid-search-hyperparameters-for-pytorch-models/>)

IEEE paper mentions using cv of 10, so that is what I'll use..

GridSearch

exhaustively searches through all possible combinations of hyperparameters during training the phase. Before we proceed further, we shall cover another cross-validation (CV) methods since tuning hyperparameters via grid search is usually cross-validated to avoid overfitting. Hence, For accelerating the running GridSearchCV we set: n_splits=3, n_jobs=2

```
In [114]: 1 #Grid Search for the below parameters
2 from sklearn.model_selection import GridSearchCV
3 params={
4     'module__dropout':[0.5,0.1],
5     'module__weight_constraint': [1.0, 2.0, 3.0, 4.0, 5.0],
6     'lr':[0.01,0.05,0.1],
7     'max_epochs':[50,100],
8     'batch_size':[50,100],
9     'optimizer__weight_decay':[0.01,0.5]
10 }
11 gs=GridSearchCV(net,params,cv=10, scoring=None,n_jobs=-1,verbose=0)
12 mlp_model=gs.fit(X_train.float(), y_train)
13 print(gs.best_score_,gs.best_params_)
```

epoch	train_loss	valid_acc	valid_loss	dur
1	2.5646	0.7765	0.7801	0.0502
2	1.0940	0.8676	0.4157	0.0608
3	0.8985	0.9265	0.4335	0.0507
4	0.8068	0.9265	0.3751	0.0512
5	0.7734	0.9324	0.3520	0.0480
6	0.7581	0.9294	0.3460	0.0403
7	0.7562	0.9294	0.3599	0.0358
8	0.7974	0.9324	0.3516	0.0500
9	0.7312	0.9294	0.3446	0.0474
10	0.7751	0.9294	0.3477	0.0364
11	0.7001	0.9294	0.3489	0.0435
12	0.7729	0.9294	0.3558	0.0387
13	0.6475	0.9324	0.3214	0.0470
14	0.7456	0.9324	0.3501	0.0426
15	0.6802	0.9324	0.3144	0.0443
16	0.8304	0.9294	0.3613	0.0374
17	0.6938	0.9324	0.3162	0.0480
18	0.6901	0.9324	0.3334	0.0596
19	0.6848	0.9353	0.3372	0.0544
20	0.7619	0.9324	0.3481	0.0377
21	0.7358	0.9324	0.3175	0.0436
22	0.7513	0.9294	0.3473	0.0437
23	0.7001	0.9324	0.3170	0.0498
24	0.7885	0.9324	0.3658	0.0413
25	0.7584	0.9324	0.3137	0.0430
26	0.7416	0.9324	0.3274	0.0524
27	0.7244	0.9324	0.3183	0.0488
28	0.7660	0.9324	0.3634	0.0508
29	0.6728	0.9324	0.3169	0.0359
30	0.7246	0.9324	0.3582	0.0369
31	0.7419	0.9324	0.3399	0.0544
32	0.6886	0.9294	0.3136	0.0362
33	0.7114	0.9324	0.3373	0.0377
34	0.7221	0.9294	0.3247	0.0411
35	0.7524	0.9324	0.3215	0.0383
36	0.7605	0.9324	0.3423	0.0411
37	0.7222	0.9324	0.3144	0.0362
38	0.7911	0.9294	0.3402	0.0456
39	0.6681	0.9353	0.3321	0.0410
40	0.7178	0.9324	0.3325	0.0396
41	0.6439	0.9294	0.3150	0.0402

Stopping since valid_loss has not improved in the last 10 epochs.

0.9288235294117648 {'batch_size': 100, 'lr': 0.01, 'max_epochs': 100, 'module__dropout': 0.1, 'module__weight_constraint': 2.0, 'optimizer__weight_decay': 0.01}

4.9 results: Severally overfitting

Stopping since valid_loss has not improved in the last 10 epochs. 0.9888888888888889 {'batch_size': 50, 'lr': 0.05, 'max_epochs': 100, 'module__dropout': 0.1, 'module__weight_constraint': 3.0, 'optimizer__weight_decay': 0.01} Stopping since valid_loss has not improved in the last 10 epochs.

Continue working through fetal class kaggle. almost have it figured out!. <https://discuss.pytorch.org/t/runtimeerror-mat1-and-mat2-must-have-the-same-dtype/166759> (<https://discuss.pytorch.org/t/runtimeerror-mat1-and-mat2-must-have-the-same-dtype/166759>) [pytorch_most_common_errors](#)

```
In [115]: 1 print(gs.best_score_,gs.best_params_)
```

0.9288235294117648 {'batch_size': 100, 'lr': 0.01, 'max_epochs': 100, 'module__dropout': 0.1, 'module__weight_constraint': 2.0, 'optimizer__weight_decay': 0.01}

```
In [ ]: 1 #mlp_model = gs.fit(X_train.float(), y_train)
```

Prediction

```
In [116]: 1 from sklearn.metrics import accuracy_score
```

```
In [117]: 1 #train
2 y_pred_train = mlp_model.predict(X_train.float())
3 #test
4 y_pred_test = mlp_model.predict(X_test.float())
```

Is the next step necessary?

```
In [118]: 1 #getting the recall_score on train
2 from sklearn.metrics import accuracy_score
3 from sklearn.metrics import recall_score
4
5 print("recall_score:", recall_score(y_train, y_pred_train, average='macro'))
6 print("accuracy_score:", accuracy_score(y_train, y_pred_train))
7 from sklearn.metrics import classification_report, confusion_matrix
8 confusion_matrix(y_train, y_pred_train)
```

```
recall_score: 0.8105692357748753
accuracy_score: 0.9294117647058824
```

```
Out[118]: array([[1300, 12, 11],
 [ 46, 188, 2],
 [ 6, 43, 92]], dtype=int64)
```

getting the recall_score on test

```
In [119]: 1 #getting the recall_score on test
2 from sklearn.metrics import accuracy_score
3 from sklearn.metrics import recall_score
4
5 print("recall_score:", recall_score(y_test, y_pred_test, average='macro'))
6 print("accuracy_score:", accuracy_score(y_test, y_pred_test))
7 from sklearn.metrics import classification_report, confusion_matrix
8 confusion_matrix(y_test, y_pred_test)
```

```
recall_score: 0.7352543345294009
accuracy_score: 0.9061032863849765
```

```
Out[119]: array([[327, 3, 2],
 [ 18, 40, 1],
 [ 1, 15, 19]], dtype=int64)
```


```
In [120]: 1 #Binarize labels in a one-vs-all fashion to use the Y_test_Roc values could be used while plotting the ROC curves
2 from sklearn.preprocessing import label_binarize
3 Y_test_ROC = label_binarize(y_test, classes=[0, 1, 2])
4 print(Y_test_ROC)
```

```
[[1 0 0]
 [1 0 0]
 [1 0 0]
 ...
 [1 0 0]
 [0 1 0]
 [1 0 0]]
```

```
In [121]: 1 report = classification_report(y_test, y_pred_test)
          2 print(report)
```

	precision	recall	f1-score	support
0	0.95	0.98	0.96	332
1	0.69	0.68	0.68	59
2	0.86	0.54	0.67	35
accuracy			0.91	426
macro avg	0.83	0.74	0.77	426
weighted avg	0.90	0.91	0.90	426

```
In [122]: 1 import seaborn as sns
          2
          3 plt.figure(figsize=(10,6))
          4 fx=sns.heatmap(confusion_matrix(y_test,y_pred_test), annot=True, fmt=".2f",cmap="GnBu")
          5 fx.set_title('Confusion Matrix MLP\n');
          6 fx.set_xlabel('\n Predicted Values\n');
          7 fx.set_ylabel('Actual Values\n');
          8 fx.xaxis.set_ticklabels(['normal', 'suspect', 'pathological'])
          9 fx.yaxis.set_ticklabels(['normal', 'suspect', 'pathological'])
         10 plt.show()
```



```
In [123]: 1 from sklearn.preprocessing import LabelBinarizer
          2
          3 label_binarizer = LabelBinarizer().fit(y_train)
          4 y_onehot_test = label_binarizer.transform(y_test)
          5 y_onehot_test.shape # (n_samples, n_classes)
```

Out[123]: (426, 3)

```
In [124]: 1 class_of_interest = 0
          2 class_id = np.flatnonzero(label_binarizer.classes_ == class_of_interest)[0]
          3 class_id
```

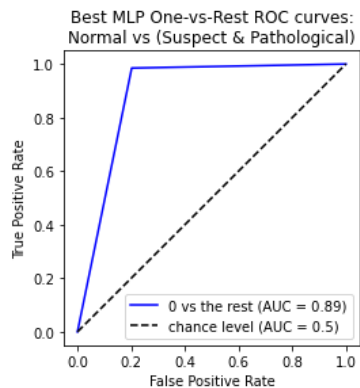
Out[124]: 0

```
In [125]: 1 predBi = label_binarizer.transform(y_pred_test)
```

```

In [126]: 1 import matplotlib.pyplot as plt
2 from sklearn.metrics import RocCurveDisplay
3
4 RocCurveDisplay.from_predictions(
5     y_onehot_test[:, class_id],
6     predBi[:, class_id],
7     name=f"{class_of_interest} vs the rest",
8     color="blue",
9 )
10 plt.plot([0, 1], [0, 1], "k--", label="chance level (AUC = 0.5)")
11 plt.axis("square")
12 plt.xlabel("False Positive Rate")
13 plt.ylabel("True Positive Rate")
14 plt.title("Best MLP One-vs-Rest ROC curves:\nNormal vs (Suspect & Pathological)")
15 plt.legend()
16 plt.show()

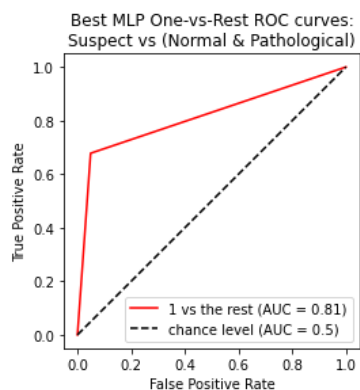
```



```

In [127]: 1 class_of_interest = 1
2 class_id = np.flatnonzero(label_binarizer.classes_ == class_of_interest)[0]
3 class_id
4
5 RocCurveDisplay.from_predictions(
6     y_onehot_test[:, class_id],
7     predBi[:, class_id],
8     name=f"{class_of_interest} vs the rest",
9     color="red",
10 )
11 plt.plot([0, 1], [0, 1], "k--", label="chance level (AUC = 0.5)")
12 plt.axis("square")
13 plt.xlabel("False Positive Rate")
14 plt.ylabel("True Positive Rate")
15 plt.title("Best MLP One-vs-Rest ROC curves:\nSuspect vs (Normal & Pathological)")
16 plt.legend()
17 plt.show()

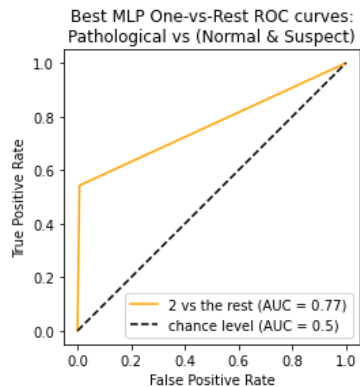
```



```

In [128]: 1 class_of_interest = 2
2 class_id = np.flatnonzero(label_binarizer.classes_ == class_of_interest)[0]
3 class_id
4
5 RocCurveDisplay.from_predictions(
6     y_onehot_test[:, class_id],
7     predBi[:, class_id],
8     name=f"{class_of_interest} vs the rest",
9     color="orange",
10 )
11 plt.plot([0, 1], [0, 1], "k--", label="chance level (AUC = 0.5)")
12 plt.axis("square")
13 plt.xlabel("False Positive Rate")
14 plt.ylabel("True Positive Rate")
15 plt.title("Best MLP One-vs-Rest ROC curves:\nPathological vs (Normal & Suspect)")
16 plt.legend()
17 plt.show()

```



Classification Report: Report which includes Precision, Recall and F1-Score. Precision - Precision is the ratio of correctly predicted positive observations to the total predicted positive observations.

Precision = $TP / (TP + FP)$

Recall (Sensitivity) - Recall is the ratio of correctly predicted positive observations to the all observations in actual class - yes.

Recall = $TP / (TP + FN)$

F1 score - F1 Score is the weighted average of Precision and Recall. Therefore, this score takes both false positives and false negatives into account. Intuitively it is not as easy to understand as accuracy, but F1 is usually more useful than accuracy, especially if you have an uneven class distribution. Accuracy works best if false positives and false negatives have similar cost. If the cost of false positives and false negatives are very different, it's better to look at both Precision and Recall.

SVM

```

In [129]: 1 from sklearn import svm

```

```

In [130]: 1 svmclassifier=svm.SVC(kernel='rbf')
2 svmclassifier.fit(X_train, y_train)
3 svmpred=svmclassifier.predict(X_test)

```

```
In [131]: 1 from sklearn.model_selection import GridSearchCV
2 from sklearn.svm import SVC
3
4 #param range
5 param_grid = {'C': [0.1, 1, 10, 100, 1000],
6               'gamma': [1, 0.1, 0.01, 0.001, 0.0001],
7               'kernel': ['rbf', 'linear', 'poly'],
8               'degree': [2, 3, 4]}
9
10 grid = GridSearchCV(SVC(), param_grid, scoring='recall_macro', refit = True, cv=10, verbose = 0)
11
12 # fitting the model for grid search
13 svmModel = grid.fit(X_train, y_train)
```

```
In [132]: 1 # print best parameter after tuning
2 print(grid.best_params_)
3
4 # print how our model looks after hyper-parameter tuning
5 print(grid.best_estimator_)
```


```
{'C': 1000, 'degree': 2, 'gamma': 1, 'kernel': 'rbf'}
SVC(C=1000, degree=2, gamma=1)
```

old best params: {'C': 100, 'degree': 2, 'gamma': 0.1, 'kernel': 'rbf'} SVC(C=100, degree=2, gamma=0.1)

```
In [133]: 1 grid_predictions = grid.predict(X_test)
2
3
4 # print classification report
5 print(classification_report(y_test, grid_predictions))
```

	precision	recall	f1-score	support
0	0.98	0.99	0.98	332
1	0.94	0.86	0.90	59
2	1.00	1.00	1.00	35
accuracy			0.97	426
macro avg	0.97	0.95	0.96	426
weighted avg	0.97	0.97	0.97	426

```
In [134]: 1 #show confuse
2
3 plt.figure(figsize=(10,6))
4 fx=sns.heatmap(confusion_matrix(y_test,grid_predictions), annot=True, fmt=".2f",cmap="GnBu")
5 fx.set_title('SVM Confusion Matrix \n');
6 fx.set_xlabel('\n Predicted Values\n');
7 fx.set_ylabel('True Values\n');
8 fx.xaxis.set_ticklabels(['normal','suspect','pathological'])
9 fx.yaxis.set_ticklabels(['normal','suspect','pathological'])
10 plt.show()
```



```
In [135]: 1 normalTN_SVM = ((51+0+0+35)/426)*100
2 normalTN_MLP = ((36+0+15+17)/426)*100
3
4
5 suspectTN_SVM = ((331+0+0+35)/426)*100
6 suspectTN_MLP = ((329+2+3+17)/426)*100
7
8 pathTN_SVM = ((331+1+51+8)/426)*100
9 pathTN_MLP = ((329+1+23+36)/426)*100
10
11
12 print("Normal TN SVM:", normalTN_SVM)
13 print("Normal TN MLP:", normalTN_MLP)
14
15 print("Suspect TN SVM:", suspectTN_SVM)
16 print("Suspect TN MLP:", suspectTN_MLP)
17
18 print("Pathlogic TN SVM:", pathTN_SVM)
19 print("Pathlogic TN MLP:", pathTN_MLP)
```

Normal TN SVM: 20.187793427230048
 Normal TN MLP: 15.96244131455399
 Suspect TN SVM: 85.91549295774648
 Suspect TN MLP: 82.3943661971831
 Pathlogic TN SVM: 91.78403755868545
 Pathlogic TN MLP: 91.31455399061032

```
In [136]: 1 print("recall_score:",recall_score(y_test, grid_predictions, average='macro'))
2 print("accuracy_score:",accuracy_score(y_test, grid_predictions))
```

recall_score: 0.9517902116942345
 accuracy_score: 0.9741784037558685

Export best model

```
In [137]: 1 bestMLP=mlp_model  
2 import joblib
```

```
In [138]: 1 filename = 'best_MLP_model.joblib'  
2 joblib.dump(bestMLP, filename)
```

```
Out[138]: ['best_MLP_model.joblib']
```

```
In [139]: 1 # Load the model from disk  
2 svmload=joblib.load('best_MLP_model.joblib')
```

```
In [ ]: 1
```